

A Space–Time Grid MILP Reference Model for Relay-Drone Disaster Delivery

Relay Drone Network project, version v5 (26.09.21.03)

Abstract

We give an exact mixed-integer linear programming (MILP) model for the relay-drone disaster-delivery problem solved by the reinforcement-learning agent and by the geometric relay heuristic in this project. The model discretises the continuous plane into a regular lattice augmented with the control centre and the destinations, and discretises time into periods of κ simulator steps chosen so that one period equals exactly one service time and exactly one lattice move. Relay positions are *not* given in advance: every lattice node is a candidate. Connectivity to the control centre is imposed in every period as a single-commodity flow on the graph induced by occupied nodes, which reproduces the simulator’s multi-hop range rule exactly. We state the discretisation conditions under which the model is feasible, prove the validity of the flow formulation and of the symmetry-breaking cuts, and validate every solution by replaying it in the simulator.

1 The continuous problem

A control centre (CC) sits at $p_0 \in \mathbb{R}^2$. A set D of M destinations at p_j , $j \in D$, each requires one parcel. A homogeneous fleet K of N drones starts at the CC, each loaded with Q parcels. Time advances in simulator steps of $\delta = 1$ s; a drone moves at most $v\delta$ per step. Delivering at a destination and reloading at the CC each freeze the drone for S steps. The mission ends when every destination has received its parcel; the objective is the makespan.

The defining constraint is communication. Two points can talk if they are within range R ; a drone is *connected* if a chain of drones links it to the CC. The simulator enforces this as a hard constraint at every step: a commanded displacement that would disconnect any drone is projected back onto the feasible set. Drones are therefore the only relays, and a drone acting as a relay cannot deliver at the same time. Table 1 lists the values used throughout.

Table 1: Simulator parameters (`disaster_relay_env_v5_26.09.15.22`).

Symbol	Meaning	Value
δ	simulator step	1 s
v	drone speed	13.89 m/s (50 km/h)
R	communication range	300 m
Q	parcel capacity	5
S	service time (unload at a destination, reload at the CC)	10 steps
ϱ	interaction radius for delivery and reload	25 m
—	map	1000×1000 m
—	reachability guarantee of the instance generator	$\ p_j - p_0\ \leq 0.9 NR$

2 Why the legacy model does not apply

The model in `legacy/milp/model.py` was written for an earlier and different problem. Five of its assumptions conflict with the problem above.

1. **One fixed destination per drone.** The map `k_dest_map` assigns drone k a single destination, so there is no routing decision, no parcel capacity and no reload. The present problem has $M \gg N$ destinations and capacity $Q < M$.
2. **Undefined position in transit.** Arcs carry travel times $\tau_{ij} \geq 2$ periods, and the flow balance leaves the drone at no node while it travels. A constraint that is imposed per period therefore cannot see the drone, so connectivity is enforced only at the endpoints of a flight that may last many periods.
3. **Bandwidth instead of range.** Connectivity is modelled as a data flow whose arc capacities follow a path-loss law P_0/d^α with low and high bandwidth demands. The simulator uses a binary rule: an edge exists if and only if $d \leq R$.
4. **No free relay placement.** The relay set R is empty and the node set consists of the CC and the destinations, so relays can only hover on top of a destination. In our problem a relay may need to stop anywhere in the plane.
5. **Makespan from the service indicator.** $C_k \geq t z_{kt}$ measures the last unloading period of a fixed assignment, not the completion of all M deliveries.

3 Discretisation

Time. Let a period be κ simulator steps, $\Delta = \kappa\delta$. We take $\kappa = S = 10$, so that service occupies exactly one period and the option horizon of the learned manager (20 steps) is exactly two periods. The per-period flight range is $\ell = v\Delta$.

Space. Let G be a square lattice of spacing h covering the bounding box of $\{p_0\} \cup \{p_j\}$ enlarged by h , from which we delete nodes farther than $0.9NR$ from the CC (unreachable, hence useless even as relays) and nodes within $0.35h$ of a point of interest. The node set is

$$N = \{0\} \cup D \cup G,$$

so the CC and every destination are exact nodes and the ϱ -radius interaction test is exact. Choosing

$$h = \frac{\ell}{\sqrt{2}}$$

makes all eight king moves feasible in one period, the diagonal move using the range exactly. The move set and the communication set are

$$A(i) = \{j \in N : \|p_i - p_j\| \leq \ell\} \ni i, \quad E = \{(i, j) \in N^2 : i \neq j, \|p_i - p_j\| \leq R\}.$$

Axis-aligned moves cover $h = \ell/\sqrt{2}$ rather than ℓ , so the lattice slows a drone by at most 29% in the axis directions; this is the principal source of discretisation error and it always acts against the model, never in its favour.

Resolution condition. The lattice must be fine enough to hold a relay chain. Serving a destination at distance d from the CC needs at least $\lceil d/R \rceil$ drones, and the chain has geometric slack $NR - d$. If the lattice cannot place the intermediate drones to within that slack the instance becomes *infeasible in the model although it is feasible in the plane*. On the test instance ($N = 3$, $d_{\max} = 791$ m, slack 109 m) the coarse lattice $h = 196$ m admits no chain to the farthest destination and the MILP is infeasible for every horizon, while $h = 98$ m admits one. A safe rule of thumb is

$$h \lesssim \frac{NR - d_{\max}}{N - 1}.$$

4 Model

4.1 Sets, parameters and variables

Sets: nodes N with the CC indexed 0, destinations $D \subset N$, drones K , periods $T = \{0, 1, \dots, \bar{T}\}$. Parameters: Q , ℓ , R , $A(i)$, E , and a tie-breaking weight $\varepsilon = 10^{-5}$. The horizon $\bar{T}$ is taken from the heuristic makespan with slack.

Table 2: Decision variables.

Variable	Domain	Meaning
y_{ikt}	$\{0, 1\}$	drone k occupies node i at the boundary of period t
s_{jkt}	$\{0, 1\}$	drone k delivers to destination j during period t
r_{kt}	$\{0, 1\}$	drone k reloads at the CC during period t
q_{kt}	$\{0, \dots, Q\}$	parcels on board drone k at the boundary of period t
ρ_{kt}	$[0, Q]$	parcels taken on during period t
g_{ijt}	$[0, N]$	connectivity flow on communication arc (i, j) in period t
C	$[0, \bar{T}]$	makespan, in periods

4.2 Objective

$$\min C + \varepsilon \sum_{i \in N \setminus \{0\}} \sum_{k \in K} \sum_{t \in T} y_{ikt} \quad (1)$$

The second term is a tie-break that prefers plans keeping drones at the CC when this does not delay the mission; with $\varepsilon \bar{T} N \ll 1$ it cannot change the optimal makespan.

4.3 Movement and position

$$\sum_{i \in N} y_{ikt} = 1 \quad \forall k \in K, t \in T \quad (2)$$

$$y_{i,k,t+1} \leq \sum_{j \in A(i)} y_{jkt} \quad \forall i \in N, k \in K, t < \bar{T} \quad (3)$$

$$y_{0k0} = 1 \quad \forall k \in K \quad (4)$$

Constraint (3) states that a drone may occupy i at $t+1$ only if it was at i or at a node from which i is reachable within one period. Because every move takes exactly one period, the position of every drone is defined at every period boundary — the property the legacy model lacks.

4.4 Deliveries, capacity and reload

$$\sum_{k \in K} \sum_{t < \bar{T}} s_{jkt} = 1 \quad \forall j \in D \quad (5)$$

$$s_{jkt} \leq y_{jkt}, \quad s_{jkt} \leq y_{j,k,t+1} \quad \forall j \in D, k, t \quad (6)$$

$$r_{kt} \leq y_{0kt}, \quad r_{kt} \leq y_{0,k,t+1} \quad \forall k, t \quad (7)$$

$$\sum_{j \in D} s_{jkt} + r_{kt} \leq 1 \quad \forall k, t \quad (8)$$

$$\sum_{j \in D} s_{jkt} \leq q_{kt} \quad \forall k, t \quad (9)$$

$$q_{k,t+1} = q_{kt} - \sum_{j \in D} s_{jkt} + \rho_{kt} \quad \forall k, t \quad (10)$$

$$\rho_{kt} \leq Q r_{kt}, \quad q_{k,t+1} \geq Q r_{kt} \quad \forall k, t \quad (11)$$

$$q_{k0} = Q \quad \forall k \quad (12)$$

The pair in (6) encodes the service freeze: the drone must be at the destination when it delivers and must still be there one period later, which is exactly the $S = \kappa$ steps during which the simulator blocks its motion. Constraints (11) together with $q \leq Q$ force a reload to fill the drone completely, as the simulator does.

4.5 Connectivity

Write $u_{it} = \sum_{k \in K} y_{ikt}$ for the occupancy of node i and set $u_{0t} := 1$. Every occupied node ships one unit of flow to the CC, which is the only sink:

$$\sum_{j: (i,j) \in E} g_{ijt} - \sum_{j: (j,i) \in E} g_{jit} = u_{it} \quad \forall i \in N \setminus \{0\}, t \in T \quad (13)$$

$$g_{ijt} \leq N u_{it} \quad \forall (i,j) \in E, t \in T \quad (14)$$

Proposition 1 (Correctness). *For a fixed occupancy u_t , the system (13)–(14) is feasible if and only if every occupied node is joined to node 0 by a path whose intermediate nodes are occupied and whose consecutive distances are at most R ; that is, if and only if the simulator's multi-hop connectivity test succeeds.*

Proof. ($\Leftarrow$) Send one unit from each occupied node along such a path; (14) holds because every arc used leaves an occupied node, and at most N units traverse any arc. ($\Rightarrow$) Let $i \neq 0$ be unoccupied. Then (14) forces $g_{ijt} = 0$ for all j , and (13) gives $\sum_j g_{jit} = 0$, so no flow passes through i . Hence all flow travels inside the occupied set. Node 0 is the only node without a conservation constraint, so the $\sum_i u_{it}$ units supplied must be absorbed there, and a flow decomposition yields a path from every occupied node to 0. $\square$

Proposition 2 (Redundancy of the head capacity). *Adding $g_{ijt} \leq N u_{jt}$ for $j \neq 0$ does not change the feasible set.*

Proof. Immediate from the second half of the previous proof: an unoccupied node carries no flow at all, so no arc entering it can be positive. $\square$

This halves the number of capacity rows, which dominate the model size ($|E| |T|$ rows rather than $2|E| |T|$).

4.6 Makespan and symmetry

$$C \geq t s_{jkt} \quad \forall j \in D, k \in K, t < \bar{T} \quad (15)$$

$$s_{jkt} = 0 \quad \forall k \in K, j \in D \text{ with } \text{idx}(j) < k + 1, t < \bar{T} \quad (16)$$

Proposition 3 (Validity of (16)). *The cut removes no solution value.*

Proof. Drones are identical, so any solution may be relabelled. Order the drones by the smallest index of a destination they serve, placing drones that serve none last. Under this labelling the smallest index served by drone k (1-based) is at least k , because the minima are distinct and increasing. Hence drone k serves no destination of index below k . $\square$

5 What the optimum means

Two effects separate the model from the continuous problem, and both must be stated.

Restriction. Drones are confined to lattice nodes and to period boundaries, and axis moves waste up to 29% of the range. Every grid plan maps to a physically flyable trajectory (each move is a straight segment of length at most ℓ flown within Δ), so the model explores a strict subset of the continuous plans. Its optimum is therefore an *upper* bound on the continuous optimum, not a certificate of optimality.

Sampling of the connectivity constraint. The model imposes connectivity at period boundaries. Between two boundaries the drones fly simultaneously and the chain may break transiently. We therefore do not claim feasibility: we *verify* it, by replaying the plan in the simulator with the same straight-line low-level controller used by all other methods and counting the steps in which the simulator’s projection had to cut a commanded displacement. As $\kappa \rightarrow 1$ the sampling becomes the simulator’s own test and the effect vanishes; the price is $|T| \propto \kappa^{-1}$.

Lower bound. A valid lower bound on the continuous optimum is

$$C_{\text{LB}} = \max \left\{ \frac{d_{\max}}{v}, \left\lceil \frac{M}{N} \right\rceil S \right\},$$

the flight time to the farthest destination and the service time that N parallel servers need for M deliveries. Together with the model optimum it brackets the true optimum.

6 Computational study

The instances come from the project’s own generator (`sample_instance`, seed 901) with $N = 3$ drones, so that the MILP, the heuristic and the learned policy all see the same geometry. Gurobi 13.0.2 ran with six threads on a machine that was concurrently training the reinforcement-learning agent. The lattice uses $\kappa = 10$ simulator steps per period and $h = 98.2$ m.

Two facts stand out. The incumbent is unstable: 5 runs of the same $M = 4$ model with different solver seeds returned makespans from 140 to 170 s. And the primal side degrades sharply with the number of destinations: at $M = 4$, 5 of 5 runs returned a feasible plan, at $M = 6$, 1 of 2 runs returned a feasible plan, at $M = 10$, 0 of 2 runs returned a feasible plan. The model is a reference

Table 3: Model size and search outcome. Makespan is in simulator steps (seconds). “best” is the best incumbent over the runs listed; the spread shows how much the incumbent varies between identical runs with different solver seeds.

M	$ N $	$\bar{T}$	binaries	runs	best	spread	bound
4	99	18	5913	5/5	140	140–170	81
6	98	23	7539	1/2	180	—	100
10	107	30	10941	0/2	none	—	—

Table 4: The best MILP plan replayed in the simulator, against the geometric relay rule with deadlock escape on the same instance. “cut” counts the drone-steps in which the simulator’s connectivity projection shortened a commanded displacement, the price of sampling connectivity only at period boundaries.

M	MILP	replayed	cut steps	rule	C_{LB}
4	140	138	22	145	50
6	180	176	59	185	53
10	none	—	—	244	57

for the smallest instances only. Where it does return a plan, the plan is good: replayed in the simulator it beats the geometric rule with deadlock escape (138 against 145 s at $M = 4$; 176 against 185 s at $M = 6$), and it never loses connectivity. The cut-step counts in Table 4 are the price of sampling connectivity at period boundaries: the simulator had to shorten that many commanded displacements, without ever disconnecting a drone.

7 Limitations

The model scales as $|N||K||T|$ binaries and $|E||T|$ continuous variables, with $|N| = \Theta(h^{-2})$ and $|T| = \Theta(\kappa^{-1})$, and the resolution condition ties h to the chain slack. The measured ceiling on this hardware is between six and ten destinations for three drones: with ten destinations the model has 10,587 binaries and 99,091 rows, and two runs of ninety and of forty minutes returned no feasible plan at all. The model is a reference point for small instances, not a solution method.

A certified lower bound on the continuous optimum would need the opposite discretisation: treat each lattice cell as a region rather than a point and test movement and communication with the *minimum* distance between cells, $\max\{0, \|p_i - p_j\| - r_i - r_j\}$ with $r_i = h/\sqrt{2}$ for a lattice cell and $r_i = 0$ for the control centre and the destinations. Every continuous trajectory then maps to a feasible cell schedule, so the relaxed optimum bounds the continuous one from below. Making that bound valid also requires relaxing the within-period timing of deliveries and of the service freeze, which the present model fixes to period boundaries; we leave the construction to future work rather than report a bound we cannot certify. It also inherits the two approximations of the previous section, and it assumes, as the simulator does, that service times are deterministic and that no drone fails.

